

Yuta Kato
6.814
Lab 3 Writeup
18 October 2014

I collaborated with Arjun Narayanan.

Initially, in the split leaf and split internal pages, I considered the field passed in to the method call to also be a legitimate field to “push up” in the tree. In other words, I took the median of all the tuples or entries in a leaf or internal page after I had theoretically inserted the field to insert. This made the code much more difficult, but I switched the algorithm to what it is now to actually pass the tests. I’m not sure why my original design is wrong, but I will try to ask about it.

I did not make any changes outside of BPlusTreeFile.java, except some changes to the BufferPool file that were recommended in Piazza.

I do not have missing or incomplete elements of code, but not all the tests pass. I liked the idea of the lab because I like datastructure work in general, but it was a very difficult lab to get working.

I spent 22 hours on this lab.